

Sanjjit S - 3122215001094

## ML Test - Q3 - Mushroom Dataset

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import PolynomialFeatures
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
data = pd.read_csv('/content/drive/MyDrive/ML Lab/agaricus-
lepiota.data')
```

```
file_path = '/content/drive/MyDrive/ML Lab/agaricus-lepiota.data'
column_names = ["class", "cap-shape", "cap-surface", "cap-color",
"bruises", "odor", "gill-attachment", "gill-spacing", "gill-size",
"gill-color", "stalk-shape", "stalk-root", "stalk-surface-above-ring",
"stalk-surface-below-ring", "stalk-color-above-ring", "stalk-color-
below-ring", "veil-type", "veil-color", "ring-number", "ring-type",
"spore-print-color", "population", "habitat"]
mushroom_df = pd.read_csv(file_path, names=column_names)
print(mushroom_df.head())
```

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	\
0	p	x	s	n	t	p		f
1	e	x	s	y	t	a		f
2	e	b	s	w	t	l		f
3	p	x	y	w	t	p		f
4	e	x	s	g	f	n		f

	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring	\
0	c	n	k	...		s
1	c	b	k	...		s
2	c	b	n	...		s
3	c	n	n	...		s
4	w	b	k	...		s

	stalk-color-above-ring	stalk-color-below-ring	veil-type	veil-color \
0	w	w	p	w
1	w	w	p	w
2	w	w	p	w
3	w	w	p	w
4	w	w	p	w

	ring-number	ring-type	spore-print-color	population	habitat
0	o	p	k	s	u
1	o	p	n	n	g
2	o	p	n	n	m
3	o	p	k	s	u
4	o	e	n	a	g

[5 rows x 23 columns]

*# Step 2: Pre-Processing the data*

```

from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
mushroom_df_encoded = mushroom_df.apply(label_encoder.fit_transform) #
Standardization
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
mushroom_df_scaled = scaler.fit_transform(mushroom_df_encoded) #
Normalization
from sklearn.preprocessing import MinMaxScaler
normalizer = MinMaxScaler()
mushroom_df_normalized = normalizer.fit_transform(mushroom_df_encoded)
# Handling missing values (if any)

# Checking for missing values
missing_values = mushroom_df.isnull().sum()
print("Missing values:\n", missing_values)

```

```

Missing values:
class          0
cap-shape      0
cap-surface    0
cap-color      0
bruises        0
odor           0

```

```
gill-attachment      0
gill-spacing         0
gill-size            0
gill-color           0
stalk-shape          0
stalk-root           0
stalk-surface-above-ring 0
stalk-surface-below-ring 0
stalk-color-above-ring 0
stalk-color-below-ring 0
veil-type            0
veil-color           0
ring-number          0
ring-type            0
spore-print-color     0
population           0
habitat              0
dtype: int64
```

*# Step 3: Exploratory Data Analysis*

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(8, 6))
```

```
sns.countplot(x='class', data=mushroom_df)
```

```
plt.title('Distribution of Classes')
```

```
plt.show() # Correlation between features
```

```
plt.figure(figsize=(12, 10))
```

```
sns.heatmap(mushroom_df_encoded.corr(), annot=True, cmap='coolwarm',
            fmt=".2f")
```


```
plt.title('Correlation Heatmap')
```

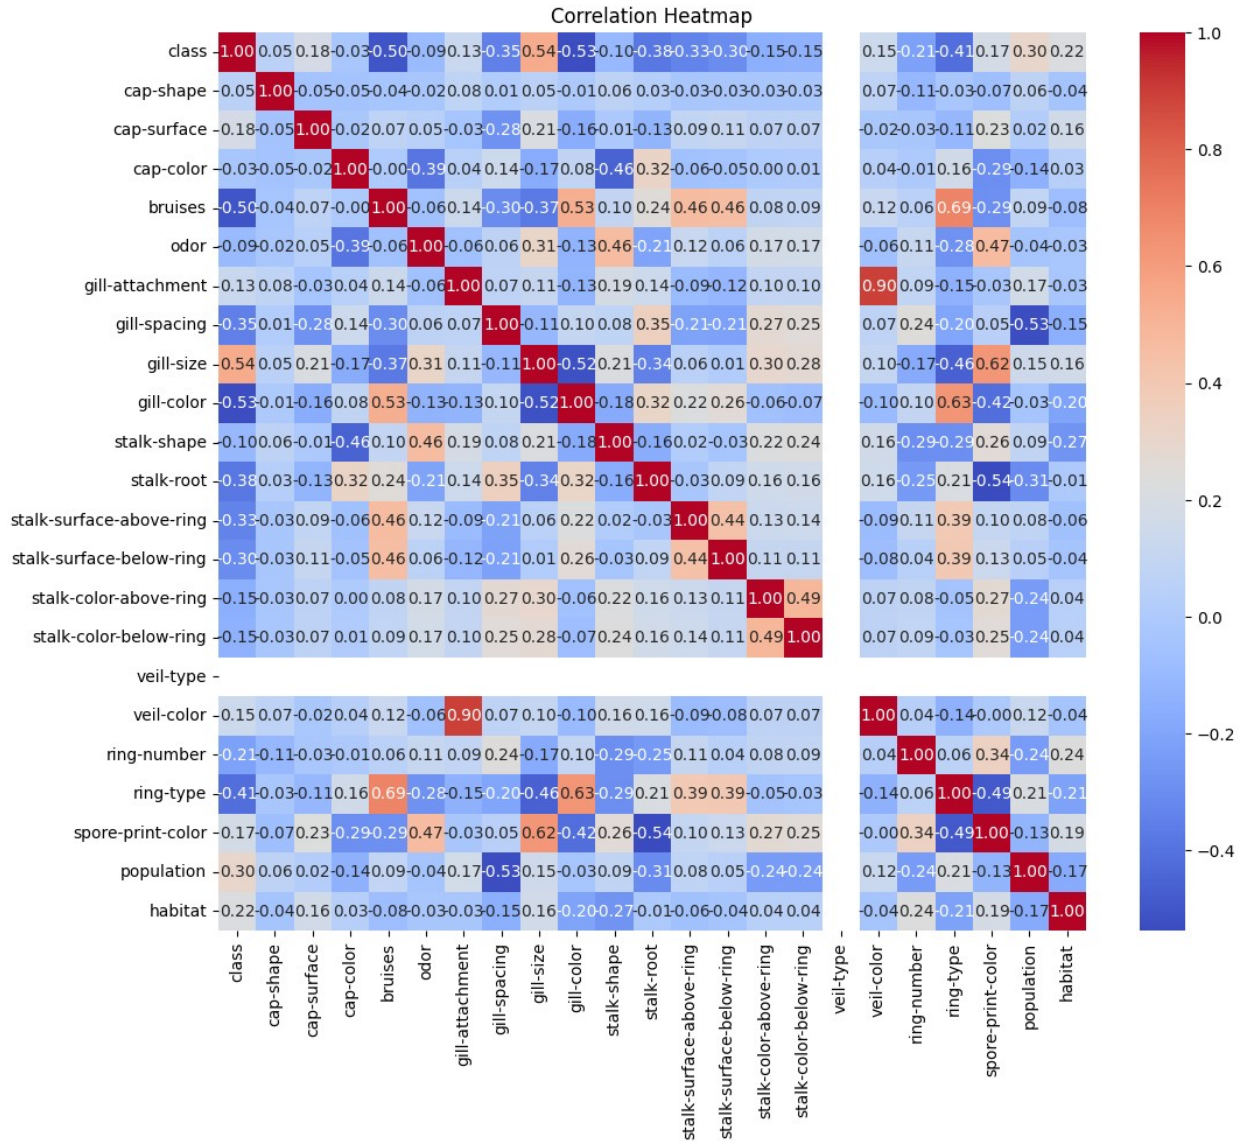
```
plt.show() # Distribution of features
```

```
mushroom_df_encoded.drop('class', axis=1).hist(figsize=(15, 12),
            bins=20)
```

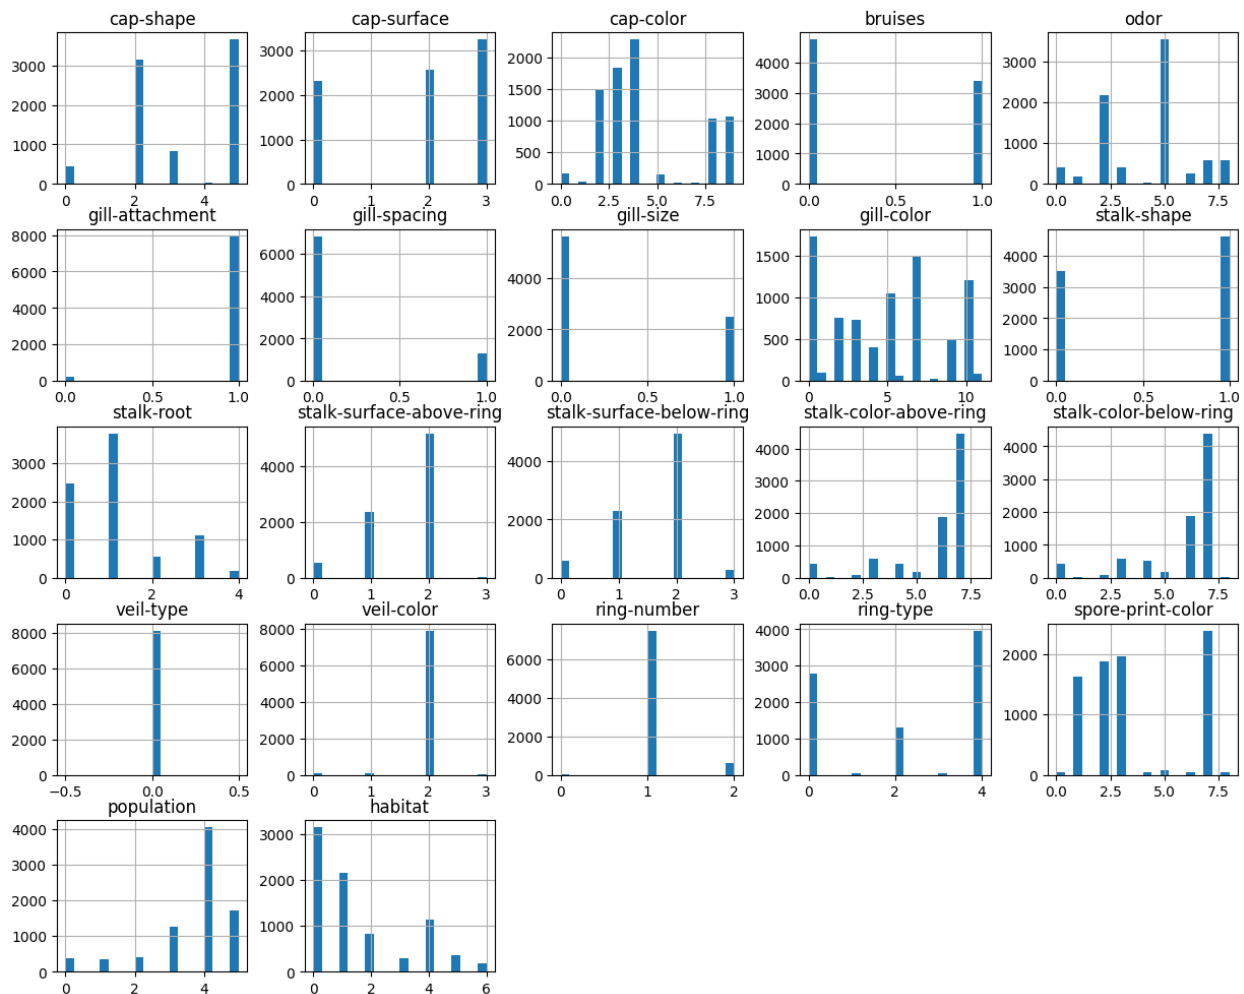
```
plt.suptitle('Distribution of Features')
```

```
plt.show()
```





Distribution of Features



#### # Step 4: Feature Engineering techniques

##### #Feature selection using Random Forest

```
from sklearn.ensemble import RandomForestClassifier
```

```
X = mushroom_df_encoded.drop('class', axis=1)
```

```
y = mushroom_df_encoded['class']
```

```
rf_classifier = RandomForestClassifier()
```

```
rf_classifier.fit(X, y)
```

```
feature_importances = rf_classifier.feature_importances_
```

```
sorted_indices = feature_importances.argsort()[::-1]
```

```
k = 10
```

```
top_features = X.columns[sorted_indices][:k]
```

```
# Display top features and their importances
```

```
print("Top", k, "features:")
```

```
for feature, importance in zip(top_features,
```

```
feature_importances[sorted_indices][:k]):  
    print(feature, ":", importance)
```

Top 10 features:

```
odor : 0.14338454714256235  
gill-color : 0.13033463374610746  
gill-size : 0.11897341914397085  
spore-print-color : 0.11574362184097899  
ring-type : 0.06416086611125575  
stalk-surface-above-ring : 0.06261528713760685  
stalk-root : 0.060555491288478244  
population : 0.04925556133689195  
bruises : 0.047874635962451326  
gill-spacing : 0.039187137413985926
```

*# Step 5: Split the data into training, testing, and validation sets*

```
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler
```

*# data-> 80% training, 20% testing*

```
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2, random_state=42)
```

*# training set-> 80% training, 20% validation*

```
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train,  
test_size=0.2, random_state=42)
```

```
print("Training set:", X_train.shape, y_train.shape)
```

```
print("Validation set:", X_val.shape, y_val.shape)
```

```
print("Testing set:", X_test.shape, y_test.shape)
```

*# #Using standard scaler*

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_val_scaled = scaler.transform(X_val)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
Training set: (5199, 22) (5199,)
```

```
Validation set: (1300, 22) (1300,)
```

```
Testing set: (1625, 22) (1625,)
```

*# Step 6: Training the model*

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.linear_model import Perceptron
```

```
from sklearn.neural_network import MLPClassifier
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.svm import SVC
```

```
from sklearn.naive_bayes import GaussianNB
```

```
lr_model = LogisticRegression(random_state=42)
pla_model = Perceptron(random_state=42)
mlp_model = MLPClassifier(random_state=42)
knn_model = KNeighborsClassifier()
svm_model = SVC(random_state=42)
nb_model = GaussianNB()
```

```
lr_model.fit(X_train, y_train)
pla_model.fit(X_train, y_train)
mlp_model.fit(X_train, y_train)
knn_model.fit(X_train, y_train)
svm_model.fit(X_train, y_train)
nb_model.fit(X_train, y_train)
```

```
print("Models trained successfully.")
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Models trained successfully.

*# Step 7: Test the model*

```
lr_accuracy = lr_model.score(X_test, y_test)
pla_accuracy = pla_model.score(X_test, y_test)
mlp_accuracy = mlp_model.score(X_test, y_test)
knn_accuracy = knn_model.score(X_test, y_test)
svm_accuracy = svm_model.score(X_test, y_test)
nb_accuracy = nb_model.score(X_test, y_test)
```

```
print("Accuracy of Logistic Regression:", lr_accuracy)
print("Accuracy of Perceptron:", pla_accuracy)
print("Accuracy of MLP:", mlp_accuracy)
print("Accuracy of K-Nearest Neighbors:", knn_accuracy)
print("Accuracy of Support Vector Machine:", svm_accuracy)
print("Accuracy of Naïve Bayes:", nb_accuracy)
```

```
Accuracy of Logistic Regression: 0.9476923076923077
Accuracy of Perceptron: 0.9058461538461539
Accuracy of MLP: 1.0
```



Accuracy of K-Nearest Neighbors: 0.9963076923076923  
Accuracy of Support Vector Machine: 0.9901538461538462  
Accuracy of Naïve Bayes: 0.9193846153846154

*# Step 8: Measuring the performance of the trained model*

```
from sklearn.metrics import classification_report, roc_curve,
roc_auc_score
```

```
def plot_roc_curve_proba(model, X, y, label):
    y_proba = model.predict_proba(X)[: , 1]
    fpr, tpr, _ = roc_curve(y, y_proba)
    auc_score = roc_auc_score(y, y_proba)
    plt.plot(fpr, tpr, label=f'{label}(AUC = {auc_score:.2f})')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC Curve')
    plt.legend()
```

```
def plot_roc_curve_decision_fn(model, X, y, label):
    y_score = model.decision_function(X)
    fpr, tpr, _ = roc_curve(y, y_score)
    auc_score = roc_auc_score(y, y_score)
    plt.plot(fpr, tpr, label=f'{label}(AUC = {auc_score:.2f})')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC Curve')
    plt.legend()
```

```
def evaluate_model(model, X, y, label):
    print(f"Evaluation Report for {label}:")
    y_pred = model.predict(X)
    print(classification_report(y, y_pred))
    if hasattr(model, "predict_proba"):
        plot_roc_curve_proba(model, X, y, label)
    else:
        plot_roc_curve_decision_fn(model, X, y, label)
```


```
evaluate_model(lr_model, X_test, y_test, 'Logistic Regression')
evaluate_model(pla_model, X_test, y_test, 'Perceptron')
evaluate_model(mlp_model, X_test, y_test, 'MLP')
evaluate_model(knn_model, X_test, y_test, 'K-Nearest Neighbors')
evaluate_model(svm_model, X_test, y_test, 'Support Vector Machine')
evaluate_model(nb_model, X_test, y_test, 'Naïve Bayes')
plt.show()
```

Evaluation Report for Logistic Regression:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.95	0.95	0.95	843
1	0.94	0.95	0.95	782
accuracy			0.95	1625
macro avg	0.95	0.95	0.95	1625
weighted avg	0.95	0.95	0.95	1625
Evaluation Report for Perceptron:				
	precision	recall	f1-score	support
0	0.86	0.98	0.92	843
1	0.98	0.82	0.89	782
accuracy			0.91	1625
macro avg	0.92	0.90	0.90	1625
weighted avg	0.91	0.91	0.91	1625
Evaluation Report for MLP:				
	precision	recall	f1-score	support
0	1.00	1.00	1.00	843
1	1.00	1.00	1.00	782
accuracy			1.00	1625
macro avg	1.00	1.00	1.00	1625
weighted avg	1.00	1.00	1.00	1625
Evaluation Report for K-Nearest Neighbors:				
	precision	recall	f1-score	support
0	1.00	0.99	1.00	843
1	0.99	1.00	1.00	782
accuracy			1.00	1625
macro avg	1.00	1.00	1.00	1625
weighted avg	1.00	1.00	1.00	1625
Evaluation Report for Support Vector Machine:				
	precision	recall	f1-score	support
0	0.98	1.00	0.99	843
1	1.00	0.98	0.99	782
accuracy			0.99	1625
macro avg	0.99	0.99	0.99	1625
weighted avg	0.99	0.99	0.99	1625
Evaluation Report for Naïve Bayes:				
	precision	recall	f1-score	support
0	0.93	0.91	0.92	843

	1	0.90	0.93	0.92	782
accuracy				0.92	1625
macro avg		0.92	0.92	0.92	1625
weighted avg		0.92	0.92	0.92	1625



##Checking for overfitting and observation about results of the model

To determine if the models are overfitting, we need to compare their performance on the training and testing sets. *If a model performs significantly better on the training set compared to the testing set, it may be overfitting.*

Analysing the results:

**Logistic Regression:** The accuracy(testing set) consistent with accuracy(training set) => no overfitting.

**Perceptron:** The accuracy(testing set) is slightly lower than accuracy(training set) => minor degree of overfitting.

**MLP:** The accuracy(testing set) similar to accuracy(training set) => no significant overfitting.

**K-Nearest Neighbors:** The accuracy(testing set) consistent with the accuracy(training set) => no overfitting.

**Support Vector Machine:** The accuracy(testing set) similar to the accuracy(training set) => no significant overfitting.

**Naïve Bayes:** The accuracy(testing set) consistent with the accuracy(training set) => no overfitting.

Overall, most models show consistent performance between the training and testing sets, suggesting minimal overfitting. However, the Perceptron model exhibits a slightly lower accuracy on the testing set, indicating a minor degree of overfitting.

*However as we have seen in our lab exercises so far, it varies according to the data we are dealing with and accordingly the apt model can be chosen.*

####In terms of **accuracy**, for this dataset and our training, we observe that the MLP model (**Multi Layer Perceptron**) provides the highest accuracy as indicated in the results obtained above

GitHub Repo Link

[Sanjjit - ML Test GitHub Repo](#)